

Haripranay Peddagolla

Arlington, Texas | (469) 893-0807 | haripranayp@gmail.com | LinkedIn

SUMMARY

Data Analyst with 4+ years of experience in fintech, banking, and technology-driven B2C/DTC industries, skilled in SQL, Python, and data visualization tools including Tableau, Power BI, and Looker. Built a real-time fraud detection ML pipeline reducing fraudulent transactions by 35% and saving an estimated \$2.5M annually. Led A/B testing initiatives resulting in a 12% uplift in customer application conversion rates. Proven ability to build KPI dashboards, ensure data governance and integrity across multiple data sources, and partner with cross-functional teams including Product, Engineering, Marketing, and Customer Success to drive data-driven decision-making. Deep expertise in dbt, dimensional modeling, Snowflake, and actionable insights that scale unit economic efficiency for consumer-facing fintech platforms.

TECHNICAL SKILLS

- **Data Analysis & BI:** SQL (PostgreSQL, MySQL, Snowflake), Advanced Excel, Power BI, Tableau, Looker, Google Data Studio, Mode Analytics
- **Programming & Scripting:** Python (Pandas, NumPy, Scikit-learn, Matplotlib, Seaborn), R (tidyverse, ggplot2, dplyr), PySpark
- **Data Engineering & ETL:** Apache Airflow, dbt, Apache Spark, AWS Glue, Azure Data Factory, Talend, SSIS, Data Transformation, Dimensional Modeling
- **Databases & Data Warehousing:** Redshift, BigQuery, Snowflake, Databricks, Oracle, MongoDB
- **Statistical & Quantitative Analysis:** A/B Testing, Hypothesis Testing, Regression Modeling, Forecasting, Time Series Analysis, Statistical Inference
- **Machine Learning & Predictive Analytics:** Feature Engineering, Model Evaluation, Supervised & Unsupervised Learning, ML Pipelines (Scikit-learn, XGBoost, LightGBM)
- **Fintech & B2C Domain Expertise:** Payment Processing, Transaction Risk Analytics, Fraud Detection, Customer Segmentation, Consumer Behavior Analytics, Credit Scoring, Unit Economics, Regulatory Compliance (KYC, AML)
- **Data Governance & Integrity:** Data Quality Frameworks, Data Dictionary Management, Cross-Source Accuracy Validation, KPI Standardization
- **Software Development & Collaboration:** Git/GitHub, Jira, Confluence, Agile/Scrum, CI/CD (Jenkins, GitLab CI), API Integration (REST, GraphQL)
- **Cloud & Big Data Ecosystem:** AWS (S3, Redshift, Lambda, Athena), Azure (Synapse, Databricks), GCP (BigQuery, Dataflow)
- **Data Visualization & Storytelling:** Interactive Dashboards, KPI Tracking, Executive Reporting, Actionable Insights Presentation

EXPERIENCE

State Street
Data Analyst

Mar 2025 – Present
USA

- Engineered and deployed a real-time Fraud Detection ML Pipeline using Python (Scikit-learn, XGBoost) on AWS SageMaker, achieving a 35% reduction in fraudulent transactions and saving an estimated \$2.5M annually, directly strengthening payment processing risk and consumer financial outcomes.
- Spearheaded a comprehensive A/B Test on the client onboarding platform targeting Credit Scoring workflow changes; statistical analysis demonstrated a significant 12% increase in application conversion rate, directly improving frictionless user experience across B2C customer touchpoints.
- Analyzed large datasets to identify trends, customer behavior insights, and high-impact opportunities, delivering actionable insights to leadership on business performance, consumer behavior, and operational trends to drive strategic decision-making.
- Orchestrated a data migration and governance project for regulatory compliance, migrating 1.2 billion transaction records from Oracle to Snowflake and Redshift using dbt and dimensional modeling, ensuring data integrity across multiple data sources and reducing reporting latency by 70%.
- Constructed a unified FinTech KPI dashboard in Power BI and Tableau consolidating payment processing and risk metrics via advanced SQL queries, improving executive decision-making speed by 20% and ensuring data accuracy and governance across business functions.

- Developed ETL jobs using Apache Airflow and PySpark alongside complex SQL transformations to ingest and cleanse market data, improving data quality from 85% to 98% accuracy and providing cross-functional teams with faster, reliable access to clean data.
- Executed a predictive modeling project using Time Series Analysis on historical trading volumes, forecasting liquidity needs with 94% accuracy and helping asset managers optimize \$50M in reserves — identifying high-impact opportunities to improve unit economic efficiency.
- Championed dbt-based data transformation workflows during large-scale financial transaction migration, standardizing models across Snowflake and Redshift and significantly improving regulatory compliance reporting efficiency.

KPIT Technologies

Jul 2020 – Aug 2023

Data Analyst

India

- Designed a process mining solution analyzing developer workflows within the Agile/Scrum framework, extracting data from Jira via REST API integration and visualizing bottlenecks in Looker and Mode Analytics, reducing ticket cycle time by 18%.
- Automated software release metrics reporting by developing Python (Pandas) scripts deployed through GitLab CI/CD, reducing weekly manual reporting effort from eight hours to under thirty minutes and enabling faster, data-driven decision-making across teams.
- Improved engineering log data quality in GCP BigQuery by identifying and resolving schema issues and enforcing data integrity and governance standards, reducing data-related errors reported by the QA team by 40% within six months.
- Implemented a proactive system failure forecasting model using Supervised Learning (Scikit-learn) on infrastructure and operational data, providing engineering teams with two-hour lead time on critical alerts to improve operational performance.
- Developed a dynamic resource allocation KPI dashboard in Google Data Studio by building SQL queries against AWS Athena data and integrating with AWS Lambda, tracking cloud spend versus project velocity and reducing average monthly cloud costs by \$10,000.
- Governed the data dictionary for a core internal application, documenting 50+ key data entities in Confluence and enforcing SQL best practices, improving cross-functional report consistency from 65% to 95% and ensuring data integrity across multiple platforms.
- Partnered with Product, Engineering, and Operations stakeholders to consolidate infrastructure, operational, and development metrics into SQL-based structured reporting pipelines, enabling faster performance monitoring and scaling data-driven decision-making.

EDUCATION

University of Texas

Master of Science, Data Science

Aug 2023 – May 2025

Arlington, Texas, USA

PROJECTS

Automated Code Review System

- Built an AI-driven automated code review system using Python (Scikit-learn) and static code analysis to detect syntax, logic, and performance issues across large code repositories, applying ML pipelines to surface actionable insights for development teams.
- Analyzed historical Git commit and code quality data to train and evaluate supervised ML models, reducing manual review effort by 40% and improving software release reliability across multiple cross-functional development teams.